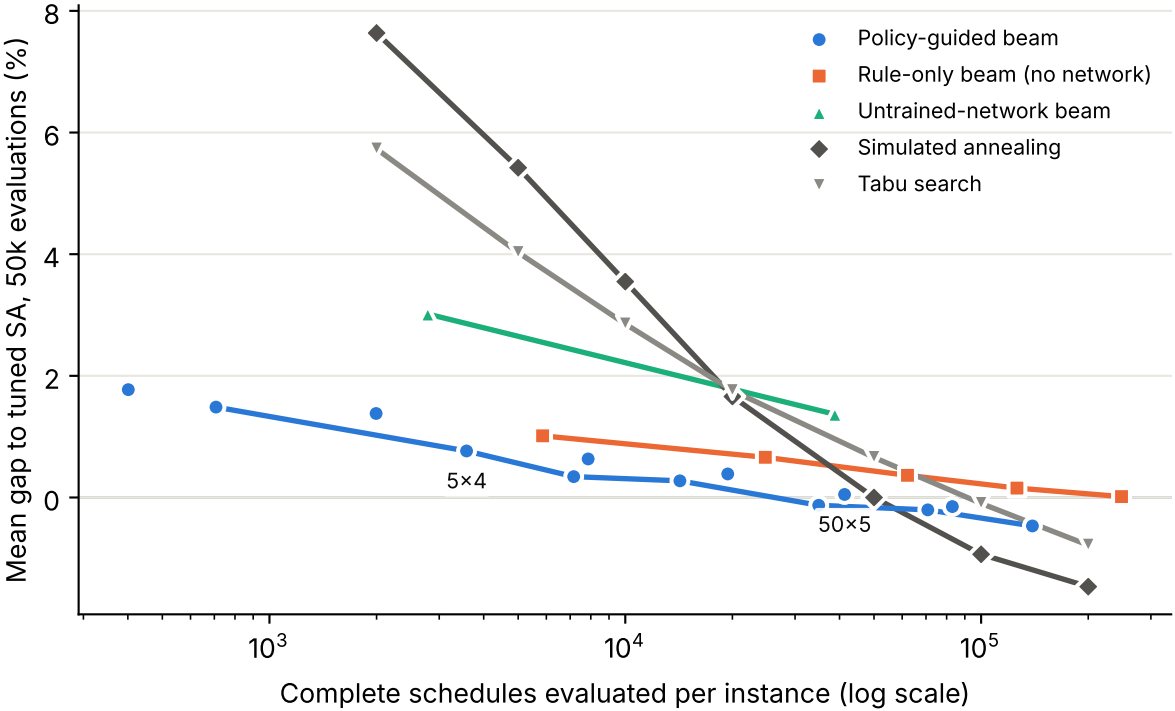
Mean gap to tuned SA, 50k evaluations (%)
8
6
4
2
0
Policy-guided beam
Rule-only beam (no network)
Untrained-network beam
Simulated annealing
Tabu search
5×4
50×5
10^3
10^4
10^5
Complete schedules evaluated per instance (log scale)