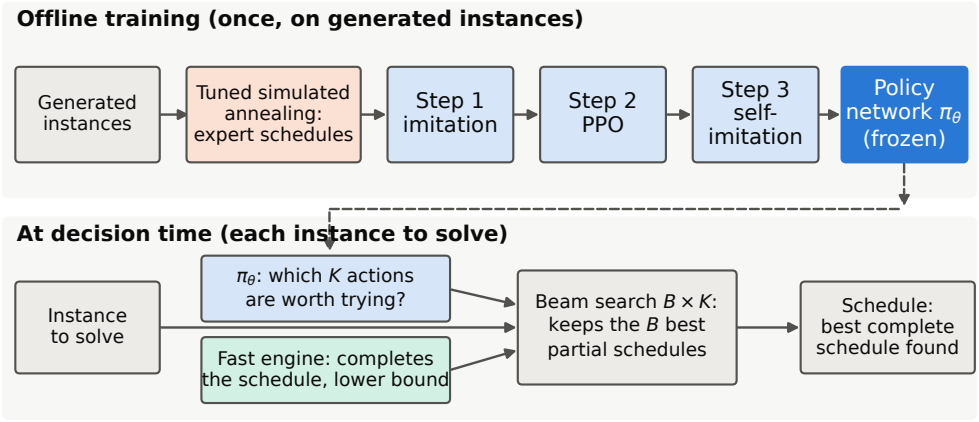
Offline training (once, on generated instances)
Generated instances
Tuned simulated annealing: expert schedules
Step 1 imitation
Step 2 PPO
Step 3 self-imitation
Policy network π_θ (frozen)
At decision time (each instance to solve)
Instance to solve
π_θ: which K actions are worth trying?
Fast engine: completes the schedule, lower bound
Beam search $B \times K$: keeps the B best partial schedules
Schedule: best complete schedule found